



Replication: Quorum based approaches + CRDTs

Philippas Tsigas
Distributed Systems, Advanced
TDA297(CTH), DIT290(GU)



What did we discuss during the previous lecture?

- ❑ **Replication for availability: Lazy Updates**
 - ❑ **Gossiping**
- ❑ **Quorum Consensus**

*“Pray with a quorum of 10
Debate with assembly of 9
Scottish dance with a collective of 8
Party with a gathering of 7
Play volleyball with a group of 6
Rank on a scale to 5
Practice music with a band of 4
Perceive in a dimension of 3
Make love with the intimacy of 2
Write poetry for an audience of 1”
— Beryl’Dov*



Conflict-free Replicated Data Types

Philippas Tsigas
Distributed Systems, Advanced
TDA297(CTH), DIT290(GU)



Motivation

- Replication and Consistency - essential features of large distributed systems such as www, p2p, and cloud computing
- Lots of replicas
 - ✓ Great for fault-tolerance and read latency
 - × Problematic when updates occur
 - Slow synchronization
 - Conflicts in case of no synchronization



Motivation

- We look for an approach that:
 - supports **Replication**
 - guarantees **Eventual Consistency**
 - is **Fast** and **Simple**
- *Conflict-free* objects = no synchronization whatsoever
- Is this practical?



Theory

Strong Eventual Consistency (SEC)

- A solution to the CAP problem
- Formal definitions
- Two sufficient conditions
- Strong equivalence between the two
- Incomparable to sequential consistency


Practice

CRDTs = Convergent or Commutative Replicated Data Types

- Counters
- Set
- Directed graph



Strong Consistency



Ideal consistency: all replicas know about the update immediately after it executes



Preclude conflicts

- Replicas update in the same total order
- Any deterministic object



Consensus

- Serialization bottleneck
- Tolerates $< n/2$ faults
- Correct, but doesn't scale



Strong Consistency

Ideal consistency: all replicas know about the update immediately after it executes




Preclude conflicts

- Replicas update in the same total order
- Any deterministic object




Consensus

- Serialization bottleneck
- Tolerates $< n/2$ faults
- Correct, but doesn't scale





Strong Consistency



Ideal consistency: all replicas know about the update immediately after it executes



Preclude conflicts

- Replicas update in the same total order
- Any deterministic object




Consensus

- Serialization bottleneck
- Tolerates $< n/2$ faults
- Correct, but doesn't scale



Strong Consistency



Ideal consistency: all replicas know about the update immediately after it executes



Preclude conflicts

- Replicas update in the same total order
- Any deterministic object




Consensus

- Serialization bottleneck
- Tolerates $< n/2$ faults
- Correct, but doesn't scale



Strong Consistency



Ideal consistency: all replicas know about the update immediately after it executes



Preclude conflicts

- Replicas update in the same total order
- Any deterministic object





Consensus

- Serialization bottleneck
- Tolerates $< n/2$ faults
- Correct, but doesn't scale





Lazy/Causal Consistency



- Update local if no causal conflict and propagate
- On conflict
 - Request
- Consensus replaced with vector clocks
 - ✓ Better performance
 - × Still complex
 - × Causal consistency





Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex




Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex




Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex




Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex




Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex





Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex





Eventual Consistency



- Update local and propagate
 - No foreground synch
 - Eventual, reliable delivery
- On conflict
 - Arbitrate
 - Roll back
- Consensus moved to background
 - ✓ Better performance
 - × Still complex




Strong Eventual Consistency



- Update local and propagate
 - No synch
 - Eventual, reliable delivery
- No conflict
 - deterministic outcome of concurrent updates
- No consensus: $\leq n-1$ faults
- Solves the CAP problem




Strong Eventual Consistency



- Update local and propagate
 - No synch
 - Eventual, reliable delivery
- No conflict
 - deterministic outcome of concurrent updates
- No consensus: $\leq n-1$ faults
- Solves the CAP problem




Strong Eventual Consistency



- Update local and propagate
 - No synch
 - Eventual, reliable delivery
- No conflict
 - deterministic outcome of concurrent updates
- No consensus: $\leq n-1$ faults
- Solves the CAP problem




Strong Eventual Consistency



- Update local and propagate
 - No synch
 - Eventual, reliable delivery
- No conflict
 - deterministic outcome of concurrent updates
- No consensus: $\leq n-1$ faults
- Solves the CAP problem



Strong Eventual Consistency



- Update local and propagate
 - No synch
 - Eventual, reliable delivery
- No conflict
 - deterministic outcome of concurrent updates
- No consensus: $\leq n-1$ faults
- Solves the CAP problem



Definition of EC

- Eventual delivery: An update delivered at some correct replica is eventually delivered to all correct replicas
- Termination: All method executions terminate
- Convergence: Correct replicas that have delivered the same updates eventually reach equivalent state
 - Doesn't preclude roll backs and reconciling



Definition of SEC

- Eventual delivery: An update delivered at some correct replica is eventually delivered to all correct replicas
- Termination: All method executions terminate
- **Strong** Convergence: Correct replicas that have delivered the same updates *have* equivalent state



CRDT design concepts

Backward-compatible with sequential datatype

If operations commute, they can be concurrent

- $add(e); rm(f) \equiv rm(f); add(e) \equiv add(e) \parallel rm(f)$

Otherwise, deterministic semantics

- Close to sequential $rm(e); add(e)$ or $add(e); rm(e)$
- Don't lose updates
- Result doesn't depend on order received
- Stable preconditions



CRDT Set design space

Many Set operations commute: $add(e) / add(f)$, $add(e) / rm(f)$, etc.

Non-commuting pair: $add(e) / rm(e)$

- sequential consistency
- last writer wins? $\{ add(e) < rmv(e) \Rightarrow e \notin S$
 $\wedge rmv(e) < add(e) \Rightarrow e \in S \}$
- error state? $\{ \perp_e \in S \}$
- add wins? $\{ e \in S \}$
- remove wins? $\{ e \notin S \}$

All deterministic, satisfy conditions



SEC is incomparable to sequential consistency

- There is a SEC object that is not sequentially-consistent:

Consider a Set CRDT S with operations $add(e)$ and $remove(e)$

- $remove(e) \rightarrow add(e) \Rightarrow e \in S$
- $add(e) \parallel remove(e') \Rightarrow e \in S \wedge e' \notin S$
- $add(e) \parallel remove(e) \Rightarrow e \in S$ (suppose add wins)

Consider the following scenario with replicas p_0, p_1, p_2 :

1. $p_0[add(e); remove(e')] \parallel p_1[add(e'); remove(e)]$
2. p_2 merges the states from p_0 and $p_1 \Rightarrow p_2: e \in S \wedge e' \in S$

The state of replica p_2 will never occur in a sequentially-consistent execution (either $remove(e)$ or $remove(e')$ must be last)



Numeric Invariants

Many applications need to enforce conditions like:

$$\text{counter} \geq K$$

E.g.:

- Number of impressions left ≥ 0
- Virtual money in a game ≥ 0



Numeric invariants

$$X \geq 0$$

Given $X = n$, there are n rights to execute *dec()*

Distribute rights among replicas

- Consume rights for *dec()*
- Create rights on *inc()*



CRDT-ish

Execute operations locally without coordination

Peer-to-peer synchronisation

Fail if not enough rights exist



Bounded Counter: API

Create(type, value);

Increment(value);

Decrement(value);

Value();

Transfer(to, qty);



Bounded Counter: increment

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Increment(10);

$$R$$

	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Increment(15);

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	0
r_3	0	0	0	0

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Increment(8);

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	8	0



Bounded Counter: increment

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	0
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	8	0



Bounded Counter: decrement

$$R$$

	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

~~decrement(15);~~

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	0
r_3	0	0	0	0

decrement(5);

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

$$R$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	8	0



Bounded Counter:

transfer

R

	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0



R

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	0
r_3	0	0	0	0

R

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0



R

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	0	0	8	0

R

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0

transfer(r_1 , 4);





Bounded Counter:

transfer

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

R_1 →

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

R_2 →

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0

R_3 →



Bounded Counter: merge

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Each replica only touches his line. Merge by taking max of each cell.

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

merge(r_1, r_2);

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0



Bounded Counter: merge

$$R_1$$

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Each replica only touches its line. Merge by taking max of each cell.

$$R_2$$

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

merge(r_1, r_2);

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0



Bounded Counter: merge

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Each replica only touches his line. Merge by taking max of each cell.

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0

merge(r_3, r_2);



Bounded Counter: merge

$$R_1$$

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	0	0	0
r_3	0	0	0	0

Each replica only touches his line. Merge by taking max of each cell.

$$R_2$$

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	15	0	5
r_3	0	0	0	0

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	4	0	8	0

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0

merge(r_3, r_2);



Bounded Counter: decrement

$$\mathbf{R}$$

	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	4	0	8	0

decrement(12);

$$\mathbf{R}$$

	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	4	0	8	0

Check local rights ≥ 12

$$\mathbf{R}$$

	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0



Bounded Counter: decrement

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	4	0	8	0

decrement(12);



R_1

R	r_1	r_2	r_3	U
r_1	10	0	0	0
r_2	0	15	0	5
r_3	4	0	8	0

R_2

R	r_1	r_2	r_3	U
r_1	0	0	0	0
r_2	0	0	0	0
r_3	4	0	8	0

Check local rights ≥ 12

$$\text{local} = R[1][1] + \sum R[i][1] - \sum R[1][j] - U[1]$$

$$14 = 10 + 4 - 0 - 0$$



Using Bounded Counter

Operation execute locally; fail if no rights available

Redistribute rights

- On-demand when needed
- Proactive

Peer-to-peer synchronization

Prototype implemented on top of Riak



Some numbers from bet365

Largest European on-line betting operator

- Bursty load: 2.5 million simultaneous users
- 1 Tb working set
- 1000s servers
- Slow users: transient inconsistency OK
- Availability, read my writes, monotonic reads
- Transparency

Before: SQLserver, doesn't scale, hours to converge

mid 2013: noSQL riak: available, siblings; ad-hoc merge (hard!)



Summary

Applications requires multiple CRDTs

- Composition (e.g. Rick Map)

Need to lower expectations...

... but still possible to enforce some invariants

- Multi-key updates: HATs
- Causality
- Numeric invariants
- General invariants: red-blue, just-right consistency



Questions?

“The problem with eventual consistency jokes is that you can't tell who doesn't get it from who hasn't gotten it.”

Next.....

Based on slides from INRIA.